

Artificial Neural Networks

Lecture 2: Linear Discriminants

Tudor Berariu
tudor.berariu@gmail.com




Faculty of Automatic Control and Computers
University Politehnica of Bucharest

Lecture : 14th of October, 2015
Last Updated: 14th of October, 2015

Today's Outline

- 1 The Classification Problem
- 2 Linear Discriminant Functions
- 3 The Perceptron



Christopher M. Bishop, *Pattern recognition and machine learning (information science and statistics)*, Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2006

Section 4.1: *Discriminant functions*

Today's Outline

- 1 The Classification Problem
- 2 Linear Discriminant Functions
- 3 The Perceptron

Today's Outline

- 1 The Classification Problem
 - Linear classifiers
 - 2 classes vs. K classes
- 2 Linear Discriminant Functions
- 3 The Perceptron

Classification

Definition

Given a data set \mathbf{X} containing N examples $\mathbf{x}^{(i)}, 1 \leq i \leq N$ in a D -dimensional space, each labeled with one of the K discrete classes $\mathcal{C}_1, \mathcal{C}_2, \dots, \mathcal{C}_K$, build a model that can **classify** new examples.

Classification

Definition

Given a data set \mathbf{X} containing N examples $\mathbf{x}^{(i)}$, $1 \leq i \leq N$ in a D -dimensional space, each labeled with one of the K discrete classes $\mathcal{C}_1, \mathcal{C}_2, \dots, \mathcal{C}_K$, build a model that can classify new examples.

- Notations used:

\mathbf{X} - the data set (collection of N examples);

$\mathbf{x}^{(i)}$ - the i^{th} example in the data set;

$$\mathbf{X} = \{(\mathbf{x}^{(1)}, t^{(1)}), \dots, (\mathbf{x}^{(N)}, t^{(N)})\}$$

$x_j^{(i)}$ - j^{th} attribute of the i^{th} example: $\mathbf{x}^{(i)} = [x_1^{(i)}, \dots, x_D^{(i)}]$

$t^{(i)}$ - the scalar label of the i^{th} example
 $(t^{(i)} \in \{\mathcal{C}_1, \dots, \mathcal{C}_K\})$;

$\mathbf{t}^{(i)}$ - the label vector of the i^{th} example

Classification

Definition


Given a data set \mathbf{X} containing N examples $\mathbf{x}^{(i)}$, $1 \leq i \leq N$ in a D -dimensional space, each labeled with one of the K discrete classes $\mathcal{C}_1, \mathcal{C}_2, \dots, \mathcal{C}_K$, build a model that can classify new examples.

- Approaches to the classification problem:
 - construct **discriminant functions**
 - compute the conditional probability $p(\mathcal{C}_k|\mathbf{x})$
 - model them directly (e.g. using parametric models)
 - learn a generative model $P(\mathbf{x}|\mathcal{C}_k)$ and use Bayes:

$$p(\mathcal{C}_k|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_k)p(\mathcal{C}_k)}{p(\mathbf{x})} \quad (1)$$

Decision boundaries

- The goal is to divide the input space into **decision regions**.
- Classes are separated by **decision boundaries** or **decision surfaces**.



Linear classification models


Definition

Models that create decision boundaries that are **linear functions of the input vector \mathbf{x}** (therefore, they are $(D - 1)$ -dimensional hyperplanes in the input space) are called **linear models**.

Linear separability

Definition

A data set which can be separated perfectly using linear decision boundaries is called **linearly separable**.



Generalized linear models

- Generalized linear models:


$$y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x} + w_0) \quad (2)$$

$f(\cdot)$ nonlinear activation function

- Decision boundaries are linear functions of \mathbf{x} even if the function $f(\cdot)$ is nonlinear.
- In today's lecture: **linear discriminants** (decision surfaces are hyperplanes).

[Bis06, pag. 180]

Single Layer Network



Today's Outline

- 1 The Classification Problem
 - Linear classifiers
 - 2 classes vs. K classes
- 2 Linear Discriminant Functions
- 3 The Perceptron

The 2-class problem

- A discriminant function:

$$y(\mathbf{x}) = \sum_{j=1}^D w_j x_j + w_0 = \mathbf{w}^T \mathbf{x} + w_0 \quad (3)$$

- $\mathbf{x} \in \mathcal{C}_1$ if $y(\mathbf{x}) \geq 0$
- $\mathbf{x} \in \mathcal{C}_2$ if $y(\mathbf{x}) < 0$
- \mathbf{w} determines the orientation of the decision boundary
- Notations:

$$\tilde{\mathbf{w}} = (w_0, \mathbf{w}^T)^T$$

$$\tilde{\mathbf{x}} = (1, \mathbf{x}^T)^T$$

$$\text{so... } y(\mathbf{x}) = \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}$$

A linear discriminant function





Figure: Illustration of the geometry of a linear discriminant function in two dimensions. The decision surface, shown in red, is perpendicular to \mathbf{x} , and its displacement from the origin is controlled by the bias parameter w_0 . Also, the signed orthogonal distance of a general point \mathbf{x} from the decision surface is given by $y(\mathbf{x})/||\mathbf{w}||$. (Taken from [Bis06, pag. 182])


K classes

- K one-versus-all classifiers




K classes

- K one-versus-all classifiers




K classes

- K one-versus-all classifiers
- $K \cdot (K - 1)/2$ one-versus-one classifiers




K classes

- K one-versus-all classifiers
- $K \cdot (K - 1)/2$ one-versus-one classifiers



K classes

- K one-versus-all classifiers
- $K \cdot (K - 1)/2$ one-versus-one classifiers
- **single classifier** with K linear functions



K classes

- K one-versus-all classifiers
- $K \cdot (K - 1)/2$ one-versus-one classifiers
- **single classifier** with K linear functions

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0} \quad (4)$$

- \mathbf{x} in class k if $y_k(\mathbf{x}) > y_j(\mathbf{x}) \quad \forall j \in \{1, \dots, K\}, j \neq k$
- decision regions are always **connected** and **convex** (proof in [Bis06, pag. 184])

Let's write some code...

- Demo: generating a synthetic linearly separable data set

Today's Outline

- 1 The Classification Problem
- 2 Linear Discriminant Functions**
- 3 The Perceptron

Today's Outline

- 1 The Classification Problem
- 2 Linear Discriminant Functions
 - Least Squares
 - Fisher's linear discriminant
- 3 The Perceptron

Linear model

- Each class ($1 \leq k \leq K$) is described by a separate linear model:

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0} \quad (5)$$

- Using the matrix notation:

$$\underbrace{\mathbf{Y}}_{N \times K} = \underbrace{\tilde{\mathbf{X}}}_{N \times (D+1)} \underbrace{\tilde{\mathbf{W}}}_{(D+1) \times K} \quad (6)$$

where $\tilde{\mathbf{W}}$ is a $(D + 1) \times K$ matrix:

$$\begin{array}{cccc} w_{10} & w_{20} & \dots & w_{K0} \\ w_{11} & w_{21} & \dots & w_{K1} \\ \vdots & \vdots & \ddots & \vdots \\ w_{1D} & w_{2D} & \dots & w_{KD} \end{array}$$

Minimizing the sum-of-squares error

- The sum-of-squares error function:

$$E(\tilde{\mathbf{W}}) = \frac{1}{2} \sum_{n=1}^N \sum_{k=1}^K (y_k(\tilde{\mathbf{x}}^{(n)}, \tilde{\mathbf{w}}_k) - t_k^{(n)})^2 \quad (7)$$

Minimizing the sum-of-squares error

- The sum-of-squares error function:

$$E(\tilde{\mathbf{W}}) = \frac{1}{2} \sum_{n=1}^N \sum_{k=1}^K (y_k(\tilde{\mathbf{x}}^{(n)}, \tilde{\mathbf{w}}_k) - t_k^{(n)})^2 \quad (7)$$

- The error is a quadratic function of the weights.
- Its derivatives w.r.t the weights will be linear functions of the weights.

Minimizing the sum-of-squares error

- The sum-of-squares error function:

$$E(\tilde{\mathbf{W}}) = \frac{1}{2} \sum_{n=1}^N \sum_{k=1}^K (y_k(\tilde{\mathbf{x}}^{(n)}, \tilde{\mathbf{w}}_k) - t_k^{(n)})^2 \quad (7)$$

- The normal equations:

$$\sum_{n=1}^N \left(\sum_{l=0}^D w_{kl} x_l^{(n)} - t_k^{(n)} \right) x_j^{(n)} = 0, \quad \forall 1 \leq j \leq K \quad (8)$$

SSE in matrix form

- The sum-of-squares error function:

$$E(\tilde{\mathbf{W}}) = \frac{1}{2} \text{Tr} \left\{ \left(\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T} \right) \left(\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T} \right)^{\top} \right\} =$$

SSE in matrix form

- The sum-of-squares error function:

$$E(\tilde{\mathbf{W}}) = \frac{1}{2} \text{Tr} \left\{ \left(\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T} \right) \left(\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T} \right)^{\text{T}} \right\} =$$
$$\frac{1}{2} \left(\text{Tr} \left\{ \tilde{\mathbf{X}}\tilde{\mathbf{W}}\tilde{\mathbf{W}}^{\text{T}}\tilde{\mathbf{X}}^{\text{T}} \right\} - \text{Tr} \left\{ \tilde{\mathbf{X}}\tilde{\mathbf{W}}\mathbf{T}^{\text{T}} \right\} - \text{Tr} \left\{ \mathbf{T}\tilde{\mathbf{W}}^{\text{T}}\tilde{\mathbf{X}}^{\text{T}} \right\} + \text{Tr} \left\{ \mathbf{T}\mathbf{T}^{\text{T}} \right\} \right)$$

SSE in matrix form

- The sum-of-squares error function:

$$E(\tilde{\mathbf{W}}) = \frac{1}{2} \text{Tr} \left\{ \left(\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T} \right) \left(\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T} \right)^{\top} \right\} =$$

$$\frac{1}{2} \left(\text{Tr} \left\{ \tilde{\mathbf{X}}\tilde{\mathbf{W}}\tilde{\mathbf{W}}^{\top}\tilde{\mathbf{X}}^{\top} \right\} - \text{Tr} \left\{ \tilde{\mathbf{X}}\tilde{\mathbf{W}}\mathbf{T}^{\top} \right\} - \text{Tr} \left\{ \mathbf{T}\tilde{\mathbf{W}}^{\top}\tilde{\mathbf{X}}^{\top} \right\} + \text{Tr} \left\{ \mathbf{T}\mathbf{T}^{\top} \right\} \right)$$

- The derivate w.r.t. $\tilde{\mathbf{W}}$:

$$\begin{aligned} \frac{\partial E(\tilde{\mathbf{W}})}{\partial \tilde{\mathbf{W}}} &= \frac{1}{2} \left(\tilde{\mathbf{X}}^{\top}\tilde{\mathbf{X}}\tilde{\mathbf{W}} + \tilde{\mathbf{X}}^{\top}\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \tilde{\mathbf{X}}^{\top}\mathbf{T} - \tilde{\mathbf{X}}^{\top}\mathbf{T} \right) \\ &= \tilde{\mathbf{X}}^{\top}\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \tilde{\mathbf{X}}^{\top}\mathbf{T} \end{aligned}$$

See matrix derivation rules here: <http://cal.cs.illinois.edu/~johannes/research/matrix%20calculus.pdf>

The pseudoinverse solution

- We minimize $E(\tilde{\mathbf{W}})$ by setting $\frac{\partial E(\tilde{\mathbf{W}})}{\partial \tilde{\mathbf{W}}}$ to zero:

$$\begin{aligned}\frac{\partial E(\tilde{\mathbf{W}})}{\partial \tilde{\mathbf{W}}} = 0 &\iff \tilde{\mathbf{X}}^T \tilde{\mathbf{X}} \tilde{\mathbf{W}} - \tilde{\mathbf{X}}^T \mathbf{T} = 0 \\ &\iff \tilde{\mathbf{X}}^T \tilde{\mathbf{X}} \tilde{\mathbf{W}} = \tilde{\mathbf{X}}^T \mathbf{T} \\ &\iff \tilde{\mathbf{W}} = (\tilde{\mathbf{X}}^T \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{X}}^T \mathbf{T} \\ &\iff \tilde{\mathbf{W}} = \tilde{\mathbf{X}}^\dagger \mathbf{T}\end{aligned}$$

Problems

- lack of robustness for outliers
- penalizes *too correct* examples
- poor results when target vectors don't have a normal distribution

Let's write some code...

- Demo: computing the weights using least-squares

Today's Outline

- 1 The Classification Problem
- 2 Linear Discriminant Functions
 - Least Squares
 - Fisher's linear discriminant
- 3 The Perceptron

The 2-class problem

- The mean vectors:

$$m_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n$$

$$m_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n$$

- **Separate the projected means**, i.e. maximize $m_2 - m_1 = \mathbf{w}^T (m_2 - m_1)$

The 2-class problem

- The mean vectors:

$$m_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n$$

$$m_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n$$

- **Separate the projected means**, i.e. maximize $m_2 - m_1 = \mathbf{w}^T (m_2 - m_1)$
- Problem: this could be achieved by increasing \mathbf{w}

An example




Figure: Projecting means

An example




Figure: Projecting means

An example




Figure: Projecting means

An example




Figure: Projecting means

An example




Figure: Projecting means

Intuition on Fisher's criterion

Definition

Maximize a function that will give a large separation between the projected class means while also giving a small variance within each class.

- the within-class variance of the projected data:

$$s_k^2 = \sum_{n \in \mathcal{C}_k} (y_n - m_k)^2 \quad (9)$$

Definition

Fisher's criterion is given by:

$$J(w) = \frac{(m_2 - m_1)^2}{s_1^2 + s_2^2} \quad (10)$$

Finding \mathbf{w}

- Rewriting Fisher's criterion:

$$J(\mathbf{w}) = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}$$

- The between-class covariance:

$$\mathbf{S}_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T$$

- The within-class covariance:

$$\mathbf{S}_W = \sum_{n \in \mathcal{C}_1} (\mathbf{x}^{(n)} - \mathbf{m}_1)(\mathbf{x}^{(n)} - \mathbf{m}_1)^T + \sum_{n \in \mathcal{C}_2} (\mathbf{x}^{(n)} - \mathbf{m}_2)(\mathbf{x}^{(n)} - \mathbf{m}_2)^T$$

- Hence, the direction of \mathbf{w}

$$\mathbf{w} \propto \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$$

Today's Outline

- 1 The Classification Problem
- 2 Linear Discriminant Functions
- 3 The Perceptron**

History

- Rosenblatt's *perceptron* (1958) represents ...
 - ... the first implemented neural network;
 - ... the first model for supervised learning.




Figure: Rosenblatt

Perceptron




Figure: Rosenblatt's *Perceptron*

General form

- two-class problems
- input space is transformed using fixed nonlinear functions $\phi(\mathbf{x})$
- uses labels $\{-1, 1\}$ for the two classes

$$y(\mathbf{x}) = f(\mathbf{w}^T \phi(\mathbf{x}))$$

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

The Perceptron Algorithm

- Goal: learn a vector \mathbf{w} such that:
 - $\mathbf{w}^T \phi(\mathbf{x}^{(n)}) > 0 \forall \mathbf{x}^{(n)} \in \mathcal{C}_1$
 - $\mathbf{w}^T \phi(\mathbf{x}^{(n)}) < 0 \forall \mathbf{x}^{(n)} \in \mathcal{C}_2$
- The *perceptron criterion* minimizes:

$$E_p(\mathbf{w}) = - \sum_{\mathcal{M}} \mathbf{w}^T \phi(\mathbf{x}^{(n)}) t^{(n)}$$

The Perceptron Algorithm

Algorithm 1 Perceptron Learning Rule

```

1: procedure THE PERCEPTRON ALGORITHM( $\mathbf{X}, \mathbf{T}, \eta$ )
2:    $\mathbf{w} \leftarrow \mathbf{0}$ 
3:   while  $E_p(\mathbf{X}, \mathbf{w}) > 0$  do
4:     choose  $\mathbf{x}^{(n)} \in \mathbf{X}$ 
5:     compute  $y(\mathbf{x}^{(n)}) = \text{sgn}(\mathbf{w}^{(t)\top} \mathbf{x}^{(n)})$ 
6:     if  $t^{(n)} \neq y(\mathbf{x}^{(n)})$  then
7:        $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t+1)} + \mathbf{x}^{(n)} t^{(n)}$ 
8:   return  $\mathbf{w}$ 

```

The Perceptron Convergence Theorem

Definition

If the data set is linearly separable, the perceptron algorithm always converges.

Demo




Figure: One update step for the perceptron

Demo




Figure: One update step for the perceptron

Demo




Figure: One update step for the perceptron

Demo




Figure: One update step for the perceptron

Demo




Figure: One update step for the perceptron

Let's write some code...

- Demo: The Perceptron

Summary

- One approach to classification is to build *discriminants*.
- The perceptron, Fisher's criterion and the LMS algorithm are all *linear discriminants* (they create *decision boundaries* that are *hyperplans*)
- The perceptron learning algorithm **always** converges if the data set **is** linearly separable.
- The perceptron learning algorithm **never** converges if the data set **is not** linearly separable.
- Minsky and Papert stated that the perceptron is incapable of generalization and its limitations also hold for the multi-layer case.
- We shall see if that is true in the following lectures.

For the exam

For the exam you should be able to ...

- ... explain what *classification*, *decision boundaries*, and *linearly separable* data sets are;
- ... explain how least squares technique, Fisher's criterion and the perceptron work (description, not formulas);

Read ...

- *Discriminant Functions* [Bis06, Section 4.1]
- *Rosenblatt's Perceptron* [Hay09, Chapter 1]

Today's Outline

4 References

References I



Christopher M. Bishop, *Pattern recognition and machine learning (information science and statistics)*, Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2006.



Simon S. Haykin, *Neural networks and learning machines*, Prentice Hall, 2009.