

Artificial Neural Networks

Lecture 3: The Backpropagation Algorithm

Tudor Berariu
tudor.berariu@gmail.com



Faculty of Automatic Control and Computers
University Politehnica of Bucharest

Lecture : 21th of October, 2015
Last Updated: 21th of October, 2015

Today's Outline

- 1 Multi-Layer Perceptrons
- 2 Forward Computation
- 3 Error functions
- 4 Error Backpropagation
- 5 Learning the weights

Today's Outline





- 1 Multi-Layer Perceptrons
- 2 Forward Computation
- 3 Error functions
- 4 Error Backpropagation
- 5 Learning the weights

Let's add more layers

- Multi-Layer Perceptrons
 - **bad name** because perceptrons have discontinuous non-linearities

Network Elements

$D = 2$ inputs




$K = 2$ outputs

Network Elements

$D = 2$ inputs

hidden layer






$K = 2$ outputs

Network Elements


$D = 2$ inputs

$$a_j^{(1)} = \sum_{i=1}^D (w_{ji}^{(1)} \cdot x_i) + w_{j0}^{(1)}$$




$K = 2$ outputs


Network Elements



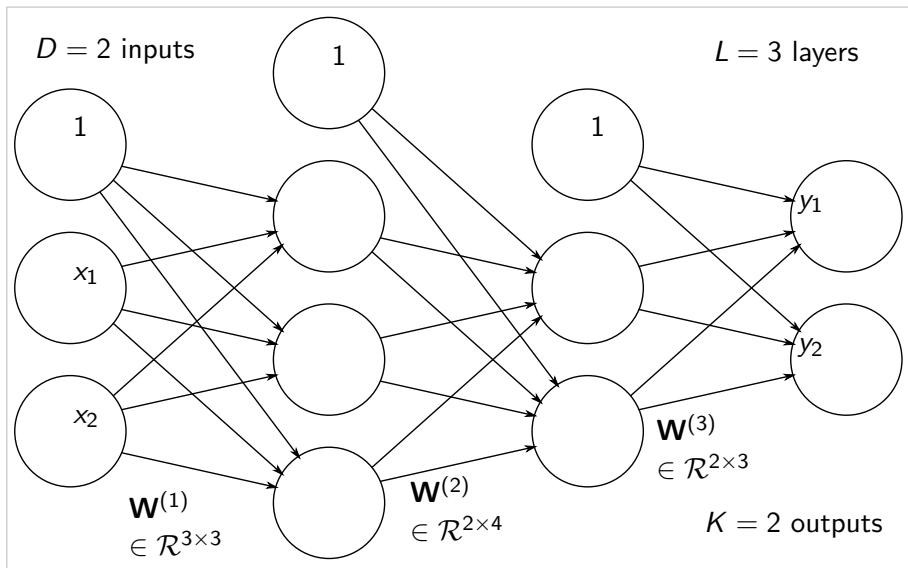
Network Elements



Network Elements



Network Elements



Notations

D - size of input space ($\mathbf{x}^{(n)} = (x_1^{(n)}, \dots, x_D^{(n)})$);

K - size of output space ($\mathbf{t}^{(n)} = (t_1^{(n)}, \dots, t_K^{(n)})$);

L number of layers (input layer is not included);

M_l number of units on layer l

- $M_0 = D$

- $M_L = K$

$\mathbf{W}^{(l)}$ weights that connect layer $l - 1$ and l

- $\mathbf{W}^{(l)} \in \mathcal{R}^{M_l \times (M_{l-1} + 1)}$

The Network function

$$y_k(\mathbf{x}, \mathbf{W}) = f_L \left(\sum_{j=1}^{M_{L-1}} w_{kj}^{(L)} z_j^{(L-1)} + w_{k0} \right)$$

The Network function

$$y_k(\mathbf{x}, \mathbf{W}) = f_L \left(\sum_{j=1}^{M_{L-1}} w_{kj}^{(L)} f_{L-1} \left(\sum_{i=1}^{M_{L-2}} w_{ji}^{(L-1)} z_i^{(L-2)} + w_{j0} \right) + w_{k0} \right)$$

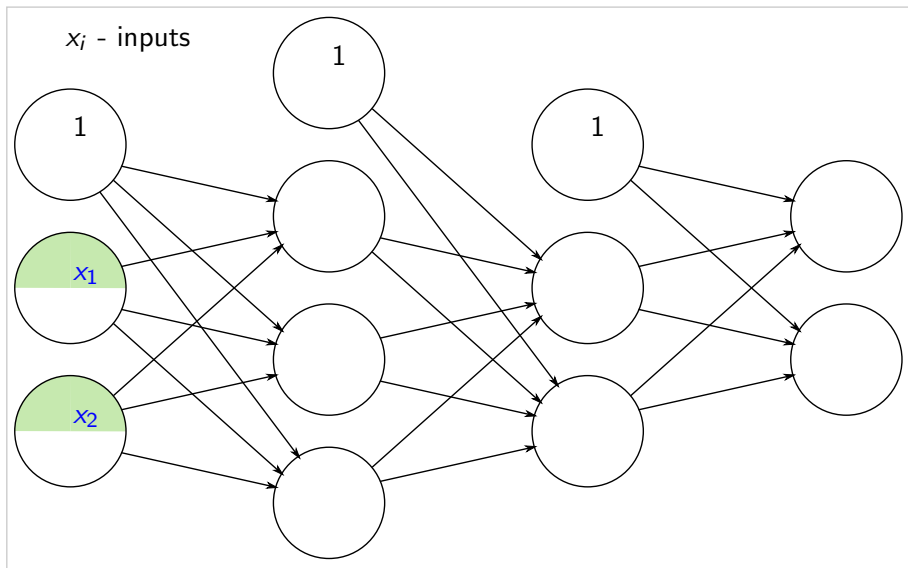
The Network function

$$y_k(\mathbf{x}, \mathbf{W}) = f_L \left(\sum_{j=1}^{M_{L-1}} w_{kj}^{(L)} f_{L-1} \left(\sum_{i=1}^{M_{L-2}} w_{ji}^{(L-1)} f_{L-2} \left(\dots \sum_{v=1}^{D+1} w_{uv}^{(1)} x_v + w_{u0}^{(1)} \dots \right) + \right. \right.$$

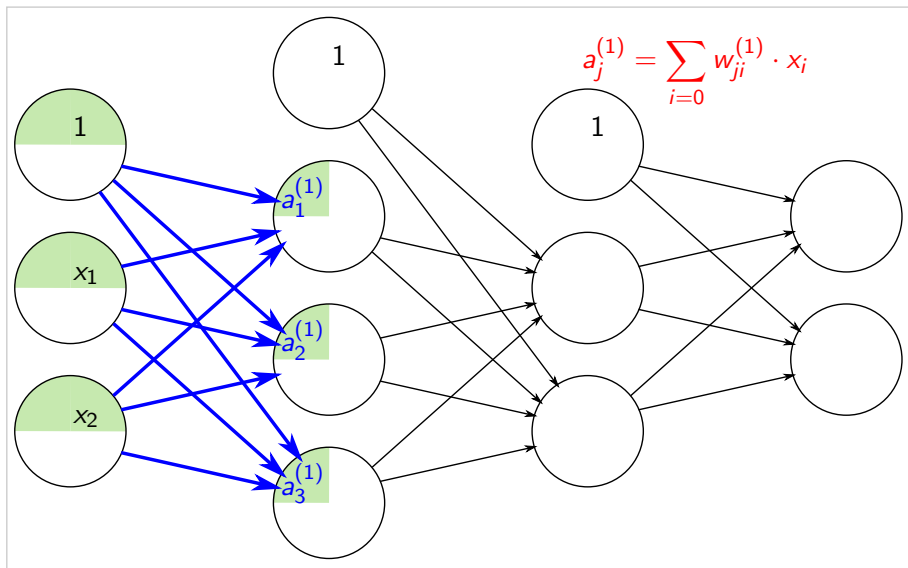
Today's Outline

- 1 Multi-Layer Perceptrons
- 2 Forward Computation**
- 3 Error functions
- 4 Error Backpropagation
- 5 Learning the weights

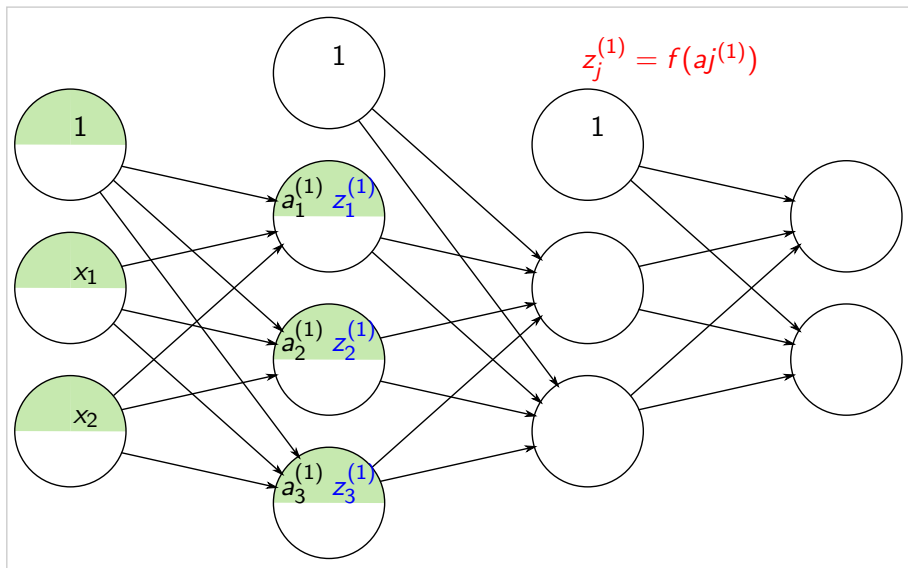
Feed-forward Computation



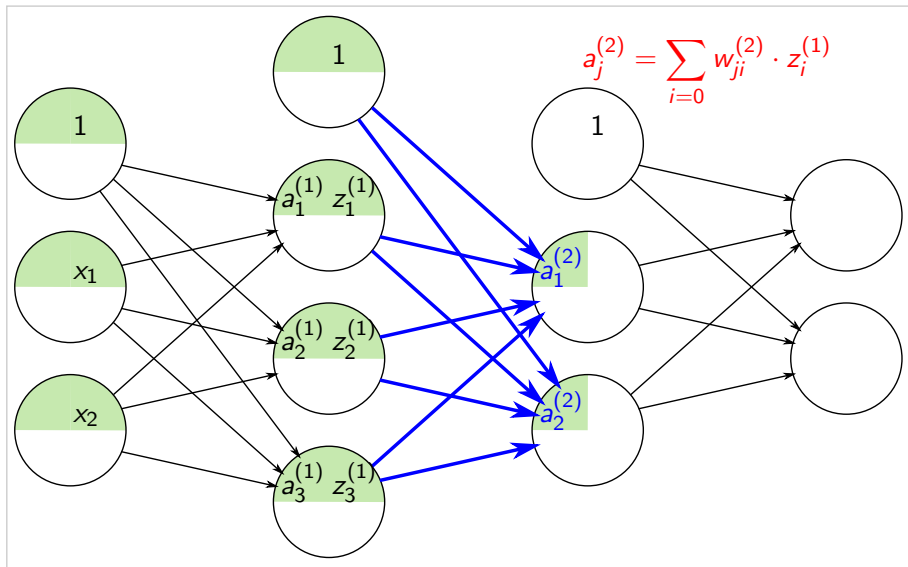
Feed-forward Computation



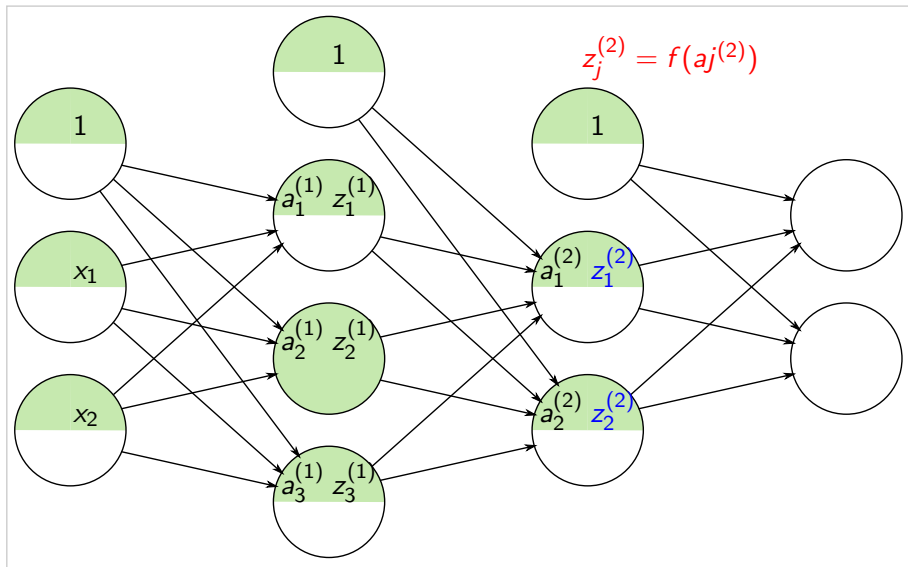
Feed-forward Computation



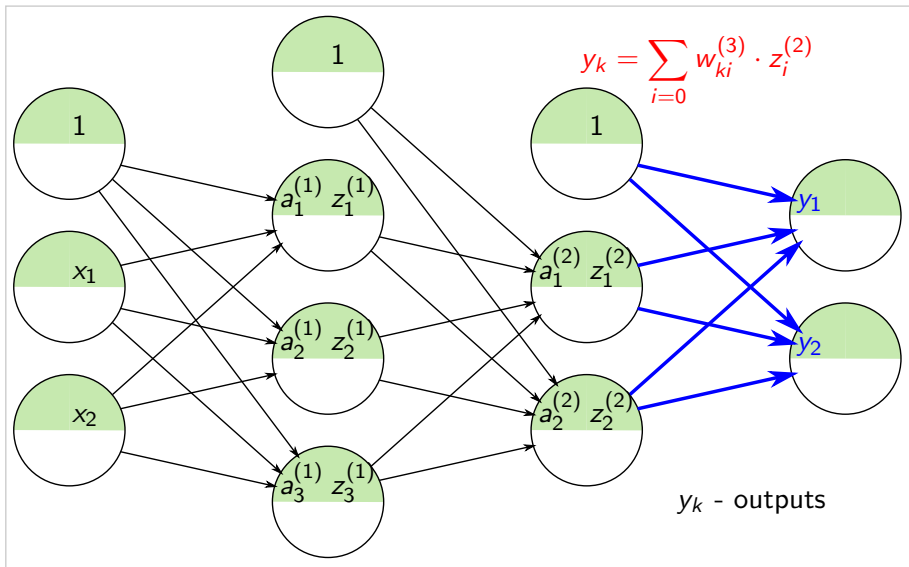
Feed-forward Computation



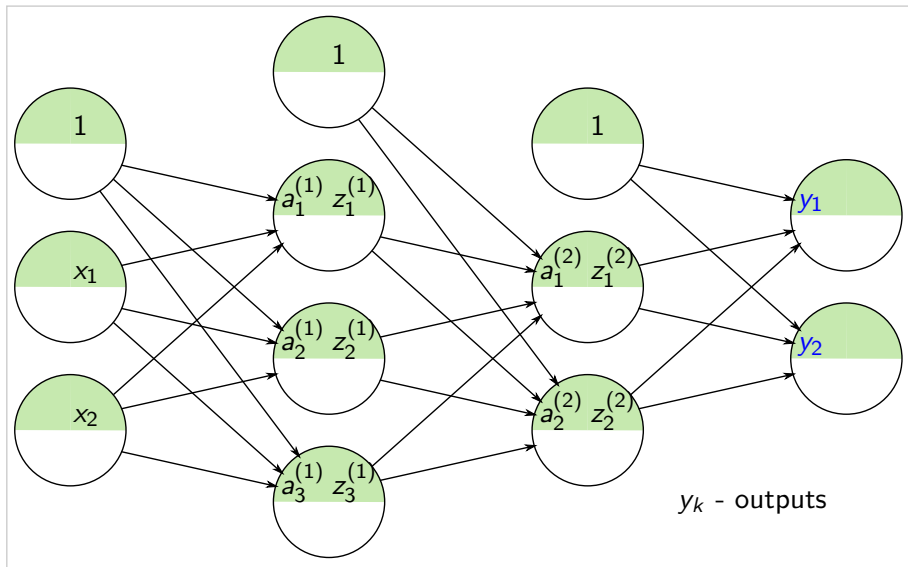
Feed-forward Computation



Feed-forward Computation



Feed-forward Computation



Today's Outline

- 1 Multi-Layer Perceptrons
- 2 Forward Computation
- 3 Error functions**
- 4 Error Backpropagation
- 5 Learning the weights

Sum of Squares for Regression

- Sum of squares error:

$$\begin{aligned} E(\mathbf{W}) &= \frac{1}{2} \sum_n^N \left\| y(\mathbf{W}, \mathbf{x}^{(n)}) - \mathbf{t}^{(n)} \right\|_2^2 \\ &= \frac{1}{2} \sum_n^N \sum_k^K (y_k(\mathbf{W}, \mathbf{x}^{(n)}) - t_k^{(n)})^2 \end{aligned}$$

Sum of Squares for Regression

- Sum of squares error:

$$\begin{aligned}
 E(\mathbf{W}) &= \frac{1}{2} \sum_n^N \left\| y(\mathbf{W}, \mathbf{x}^{(n)}) - \mathbf{t}^{(n)} \right\|_2^2 \\
 &= \frac{1}{2} \sum_n^N \sum_k^K (y_k(\mathbf{W}, \mathbf{x}^{(n)}) - t_k^{(n)})^2
 \end{aligned}$$

- its derivate w.r.t. y_k :

$$\frac{\partial E(\mathbf{W})}{\partial y_k} = y_k - t_k$$

Cross-Entropy for Classification (I)

- The network models the conditional distribution:

$$y_k(\mathbf{x}) = p(C_k|\mathbf{x})$$

- The likelihood of observing the data set (assuming i.i.d. data):

$$\begin{aligned} p(\mathbf{T}|\mathbf{X}) &= \prod_n^N p(\mathbf{t}^{(n)}|\mathbf{x}^{(n)}) \\ &= \prod_n^N \prod_k^K \left(y_k(\mathbf{x}^{(n)}) \right)^{t_k^{(n)}} \end{aligned}$$

Cross-Entropy for Classification (II)

- maximizing likelihood is the same with minimizing log-likelihood:

$$E = - \sum_n^N \sum_k^K t_k^{(n)} \log y_k(\mathbf{x}^{(n)})$$

- if softmax is used on the last layer:

$$y_k = a_k^{(L)} = \frac{e^{a_k^{(L)}}}{\sum_{k'}^K e^{a_{k'}^{(L)}}}$$

- outputs are interpreted as conditional probabilities $p(C_k|\mathbf{x})$

Cross-Entropy for Classification (III)

- the derivatives of the softmax error function

$$\frac{\partial E^{(n)}}{\partial a_k} = \sum_{k'} \frac{\partial E^{(n)}}{\partial y_{k'}} \frac{\partial y_{k'}}{\partial a_k}$$

where

$$\frac{\partial y_{k'}}{\partial a_k} = y_{k'} \delta_{kk'} - y_{k'} y_k$$

and

$$\frac{\partial E^{(n)}}{\partial y_{k'}} = -\frac{t_{k'}}{y_{k'}}$$

leads to:

$$\frac{\partial E^{(n)}}{\partial a_k} = y_k - t_k$$

Today's Outline

- 1 Multi-Layer Perceptrons
- 2 Forward Computation
- 3 Error functions
- 4 Error Backpropagation**
- 5 Learning the weights

The Training Process

- the training process for a feed-forward neural network:
 - 1 evaluation of the derivatives of the error function $E(\mathbf{w})$
 - 2 the derivatives are used to adjust the weights

The Training Process

- the training process for a feed-forward neural network:
 - ① evaluation of the derivatives of the error function $E(\mathbf{w})$ - **backpropagation**
 - ② the derivatives are used to adjust the weights - **gradient descent**

The Training Process

- the training process for a feed-forward neural network:
 - ① evaluation of the derivatives of the error function $E(\mathbf{w})$ - **backpropagation**
 - not only for the sum-of-squares error
 - ② the derivatives are used to adjust the weights - **gradient descent**
 - there are other optimization techniques beside gradient descent

The error function

Definition

$$E(\mathbf{w}) = \sum_n^N E(\mathbf{w}, \mathbf{x}^{(n)}) = \sum_n^N E_n \quad (1)$$

Assumption: i.i.d. data.

The error function

Definition

$$E(\mathbf{w}) = \sum_n^N E(\mathbf{w}, \mathbf{x}^{(n)}) = \sum_n^N E_n \quad (1)$$

Assumption: i.i.d. data.

- The goal: evaluate $\nabla E(\mathbf{w})$

The derivatives of the error

- Applying the chain rule for partial derivatives:

$$\frac{\partial E_n}{\partial w_{ji}^{(l)}} = \frac{\partial E_n}{\partial a_j^{(l)}} \cdot \frac{\partial a_j^{(l)}}{\partial w_{ji}^{(l)}} \quad (2)$$

The derivatives of the error

- Applying the chain rule for partial derivatives:

$$\frac{\partial E_n}{\partial w_{ji}^{(l)}} = \frac{\partial E_n}{\partial a_j^{(l)}} \cdot \frac{\partial a_j^{(l)}}{\partial w_{ji}^{(l)}} \quad (2)$$

- Remember that:

$$a_j^{(l)} = \sum_i w_{ji}^{(l)} \cdot z_i^{(l-1)}$$

$$\frac{\partial a_j^{(l)}}{\partial w_{ji}^{(l)}} = z_i^{(l-1)}$$

The derivatives of the error

- Applying the chain rule for partial derivatives:

$$\frac{\partial E_n}{\partial w_{ji}^{(l)}} = \frac{\partial E_n}{\partial a_j^{(l)}} \cdot \frac{\partial a_j^{(l)}}{\partial w_{ji}^{(l)}} \quad (2)$$

- Remember that:

$$a_j^{(l)} = \sum_i w_{ji}^{(l)} \cdot z_i^{(l-1)}$$

$$\frac{\partial a_j^{(l)}}{\partial w_{ji}^{(l)}} = z_i^{(l-1)}$$

- Notation (δ - errors)

$$\delta_j^{(l)} = \frac{\partial E_n}{\partial a_j^{(l)}}$$

The derivatives of the error

- Applying the chain rule for partial derivatives:

$$\frac{\partial E_n}{\partial w_{ji}^{(l)}} = \frac{\partial E_n}{\partial a_j^{(l)}} \cdot \frac{\partial a_j^{(l)}}{\partial w_{ji}^{(l)}} = \delta_j^{(l)} \cdot z_i^{(l-1)} \quad (2)$$

- Remember that:

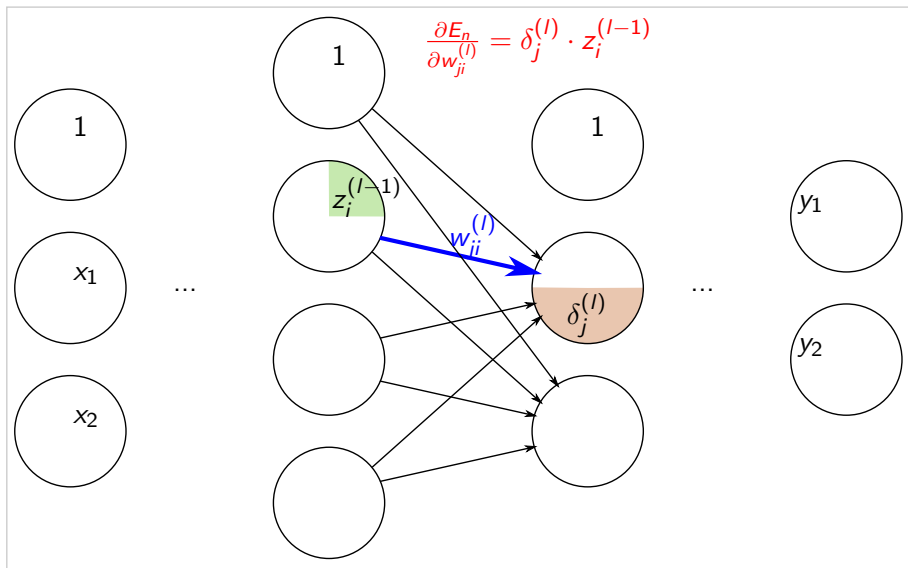
$$a_j^{(l)} = \sum_i w_{ji}^{(l)} \cdot z_i^{(l-1)}$$

$$\frac{\partial a_j^{(l)}}{\partial w_{ji}^{(l)}} = z_i^{(l-1)}$$

- Notation (δ - errors)

$$\delta_j^{(l)} = \frac{\partial E_n}{\partial a_j^{(l)}}$$

The derivatives of the error



Computing δ 's

- On the last layer ($l = L, \forall k = 1 \dots K$):

$$\delta_k^{(L)} = \frac{\partial E_n}{\partial a_k^{(L)}} = \frac{\partial E_n}{\partial y_k} \quad (3)$$

Computing δ 's

- On the last layer ($l = L, \forall k = 1 \dots K$):

$$\delta_k^{(L)} = \frac{\partial E_n}{\partial a_k^{(L)}} = \frac{\partial E_n}{\partial y_k} = y_k - t_k^{(n)} \quad (3)$$

Computing δ 's

- On the last layer ($l = L, \forall k = 1 \dots K$):

$$\delta_k^{(L)} = \frac{\partial E_n}{\partial a_k^{(L)}} = \frac{\partial E_n}{\partial y_k} = y_k - t_k^{(n)} \quad (3)$$

- On the other layers ($l < L, \forall k = 1 \dots M_l$):

$$\delta_j^{(l)} = \frac{\partial E_n}{\partial a_j^{(l)}} = \sum_k \frac{\partial E_n}{\partial a_k^{(l+1)}} \cdot \frac{\partial a_k^{(l+1)}}{\partial a_j^{(l)}} \quad (4)$$

Computing δ 's

- On the last layer ($l = L, \forall k = 1 \dots K$):

$$\delta_k^{(L)} = \frac{\partial E_n}{\partial a_k^{(L)}} = \frac{\partial E_n}{\partial y_k} = y_k - t_k^{(n)} \quad (3)$$

- On the other layers ($l < L, \forall k = 1 \dots M_l$):

$$\delta_j^{(l)} = \frac{\partial E_n}{\partial a_j^{(l)}} = \sum_k \frac{\partial E_n}{\partial a_k^{(l+1)}} \cdot \frac{\partial a_k^{(l+1)}}{\partial a_j^{(l)}} \quad (4)$$

- Remember that:

- $a_k^{(l+1)} = \sum_j w_{kj}^{(l+1)} \cdot z_j^{(l)}$
- $z_j^{(l)} = f(a_j^{(l)})$

Computing δ 's

- On the last layer ($l = L, \forall k = 1 \dots K$):

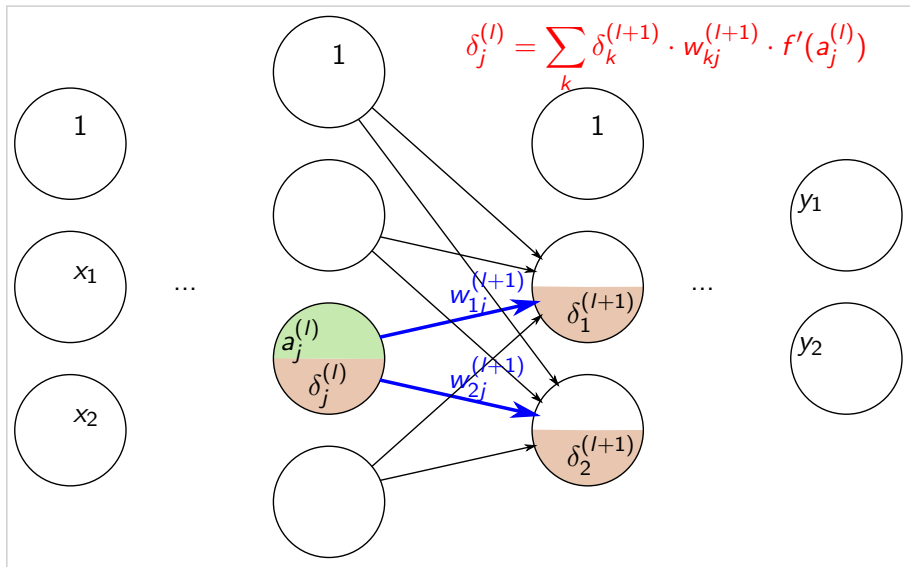
$$\delta_k^{(L)} = \frac{\partial E_n}{\partial a_k^{(L)}} = \frac{\partial E_n}{\partial y_k} = y_k - t_k^{(n)} \quad (3)$$

- On the other layers ($l < L, \forall k = 1 \dots M_l$):

$$\delta_j^{(l)} = \frac{\partial E_n}{\partial a_j^{(l)}} = \sum_k \frac{\partial E_n}{\partial a_k^{(l+1)}} \cdot \frac{\partial a_k^{(l+1)}}{\partial a_j^{(l)}} = \sum_k \delta_k^{(l+1)} \cdot w_{kj}^{(l+1)} \cdot f'(a_j^{(l)}) \quad (4)$$

- Remember that:
 - $a_k^{(l+1)} = \sum_j w_{kj}^{(l+1)} \cdot z_j^{(l)}$
 - $z_j^{(l)} = f(a_j^{(l)})$

The derivatives of the error



Error backpropagation algorithm

Algorithm 1 Error backpropagation

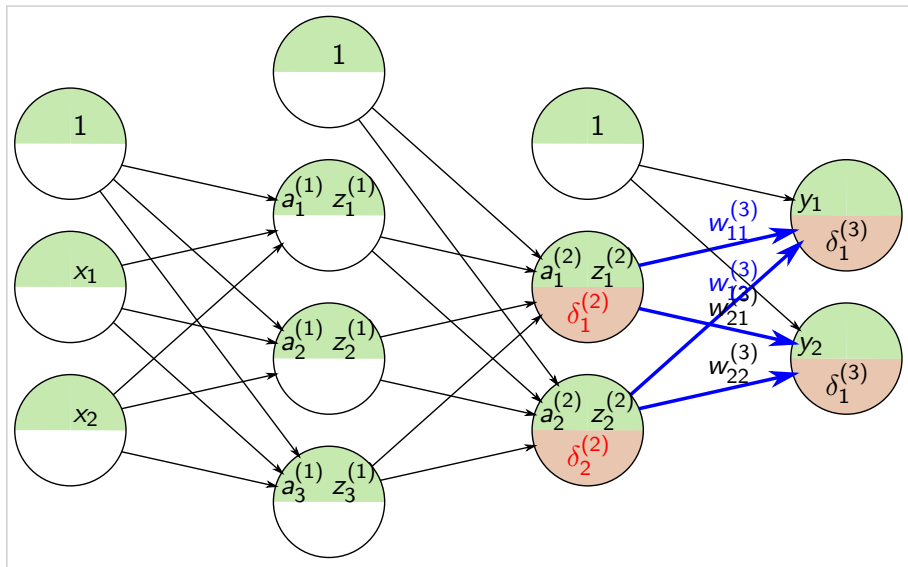
```

1: for  $k = 1 \dots K$  do
2:    $\delta_k^{(L)} \leftarrow y_k - t_k$ 
3:    $\frac{\partial E_n}{\partial w_{ki}^{(L)}} \leftarrow \delta_k^{(L)} \cdot z_i^{(L-1)}$ 
4: for  $l = (L - 1) \dots 1$  do
5:   for  $j = 1 \dots M_l$  do
6:      $\delta_j^{(l)} \leftarrow \sum_k \delta_k^{(l+1)} \cdot w_{kj}^{(l+1)} \cdot f'(a_j^{(l)})$ 
7:      $\frac{\partial E_n}{\partial w_{ji}^{(l)}} \leftarrow \delta_j^{(l)} \cdot z_i^{(l-1)}$ 

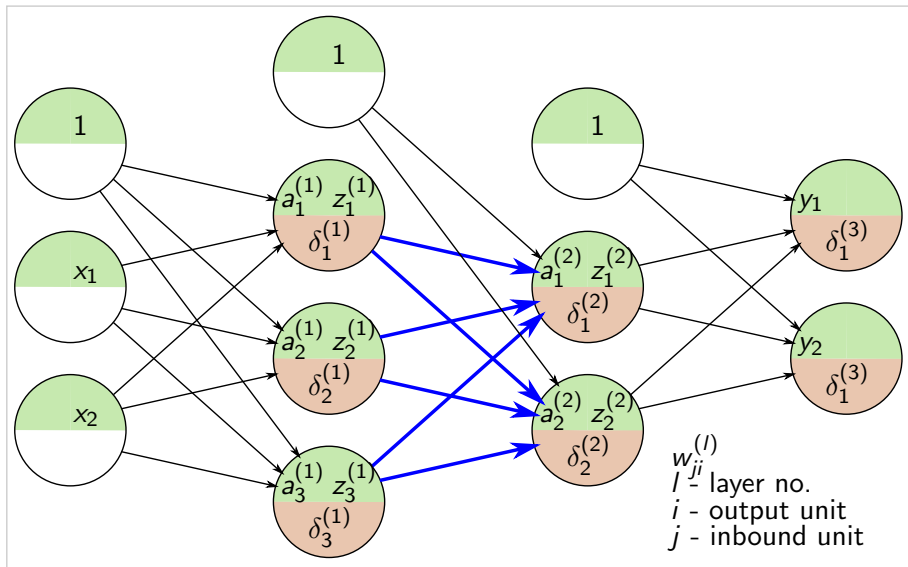
```

Error Backpropagation Example

Error Backpropagation Example



Error Backpropagation Example



Today's Outline

- 1 Multi-Layer Perceptrons
- 2 Forward Computation
- 3 Error functions
- 4 Error Backpropagation
- 5 Learning the weights

Gradient descent

- today we consider the learning rate constant
- Gradient descent algorithm:
 - until *convergence*:
 - 1 Compute $\frac{\partial E_n}{\partial w_{ji}^{(l)}}$ using backpropagation
 - 2 $\mathbf{W}^{t+1} \leftarrow \mathbf{W}^{(t)} - \eta \nabla E$

Questions


- How often do we change the weights?
- How do we prevent overfitting? Does the network generalize?

Stochastic vs. batch learning

- Stochastic - update weights after each example
 - better for on-line learning
 - better for large redundant data sets
- Batch - update weights after computing error for all training examples
 - supports parallelization
 - better estimation of the error
- Mini-batch

Learning rate

- fixed or adaptive learning rate
- are there any better directions than the steepest descent?



Prevent overfitting

- use a validation set




Figure: MSE for training and validation sets

- use regularization
 - optimize the risk: $\mathbf{R} = \alpha E(\mathbf{x}, \mathbf{W}) + (1 - \alpha)C(\mathbf{W})$
 - C - complexity penalty
 - $C = \sum_{w \in \mathbf{W}} ||w||$

Next time..

- strategies to adjust the learning rate
- regularization techniques
- better directions to move along the error surface

Summary

- Adjusting network parameters means:
 - 1 compute the derivatives of the error with respect to the parameters:
$$\frac{\partial E}{\partial w_{ij}^{(l)}}$$
 - 2 use those derivatives to optimize weights *discriminants*.
- Error backpropagation gets the gradient ∇E
- Gradient descent optimizes the weights
 - It can be done *stochastic*, *batch*, or the *mini-batch* way

For the exam

For the exam you should know ...

- ... the backpropagation method to compute $\frac{\partial E}{\partial w_{ij}^{(l)}}$
- ... gradient descent optimizations

Read ...

- *Multilayer Perceptrons* [Hay09, Chapter 4]

Today's Outline

6 References

References I



Simon S. Haykin, *Neural networks and learning machines*, Prentice Hall, 2009.